

DEPARTMENT OF ELECTRICAL AND COMPUTER ENGINEERING
UNIVERSITY OF BRITISH COLUMBIA
CPEN 211 Computer Systems I, Fall 2023

Lab 5: Datapath of the “Simple RISC Machine”

Week of Nov 6 to 10 (your code must be submitted by 9:59 PM the evening before your lab session)

1 Introduction

Starting with this lab you will build a computer that implements a “Reduced Instruction Set Computer” (RISC). The ARMv7 that we learned about in lectures and used in Lab 4 is a RISC “instruction set architecture” (ISA). In Lab 5 to 7 you will build a RISC processor that, while much simpler than ARMv7, will give you a strong grasp of how computers actually work.

The “Simple RISC Machine” executes programs written using a *very* small set of instructions. Section 8 at the end of this handout provides a complete list of the instructions your computer will execute at end of Lab 7. This set of instructions is “Turing-complete” meaning your Simple RISC Machine computer could in principle run any program given enough time and enough memory.

To start, consider the following line of C code:

```
f = (g + h) - (i + j);
```

Assuming f, g, h, i, j are in registers R0 through R4, this C code can be implemented using the Simple RISC Machine (SRM) instructions:

```
ADD R5, R1, R2; // t1 = g + h
ADD R6, R3, R4; // t2 = i + j
SUB R0, R5, R6; // f = t1 - t2
```

Where we used R5 and R6 as temporary registers.

To implement a computer that can execute these instructions we will need a block of hardware that can add and subtract numbers. We also need some hardware that can store the numbers before and after the addition and subtraction operations. The following section describes one hardware design for a computer that can execute ADD and SUB instructions like those above.

1.1 Overview of the Simple RISC Machine Datapath

To execute the instructions above along with some others described later, in Lab 5, you implement the computer *datapath* shown in Figure 1. This section explains the different elements in Figure 1. While reading this section try to focus on understanding what each individual block does. In Section 2 we will see an example showing how the blocks work together.

The datapath consists of one register file ① containing 8 registers, each holding 16-bits; three multiplexers ⑥⑦⑨; three 16-bit registers with load enable ③④⑤; a 1-bit register with load enable ⑩; a shifter unit ⑧; and an arithmetic logic unit (ALU) ②. Below we describe each component in detail.

1.1.1 Register File

The Simple RISC Machine employs the hardware structure known as a *register file* ①, shown in Figure 1. The register file is a small read-write memory containing eight locations each storing 16-bits. The number in the register name, e.g., “5” in R5, is used as the address to select one of the registers to read or write. To specify one of the eight locations requires only $\log_2(8) = 3$ -bits (even though each location can hold 16-bits) for the address.

An *individual* 16-bit register inside the register file is built using what is sometimes referred to as a *register with load enable*. Implement a register with load enable by writing (System)Verilog for the circuit in Figure 2. In Figure 2, when load is 0, out is passed through the top input of the multiplexer into the D input of the set of 16 flip-flops, which together form a 16-bit register. Thus, when load is 0 and there is a rising edge of the clock, the value in the 16-bit register does *not* change. Conversely, when load is 1 the

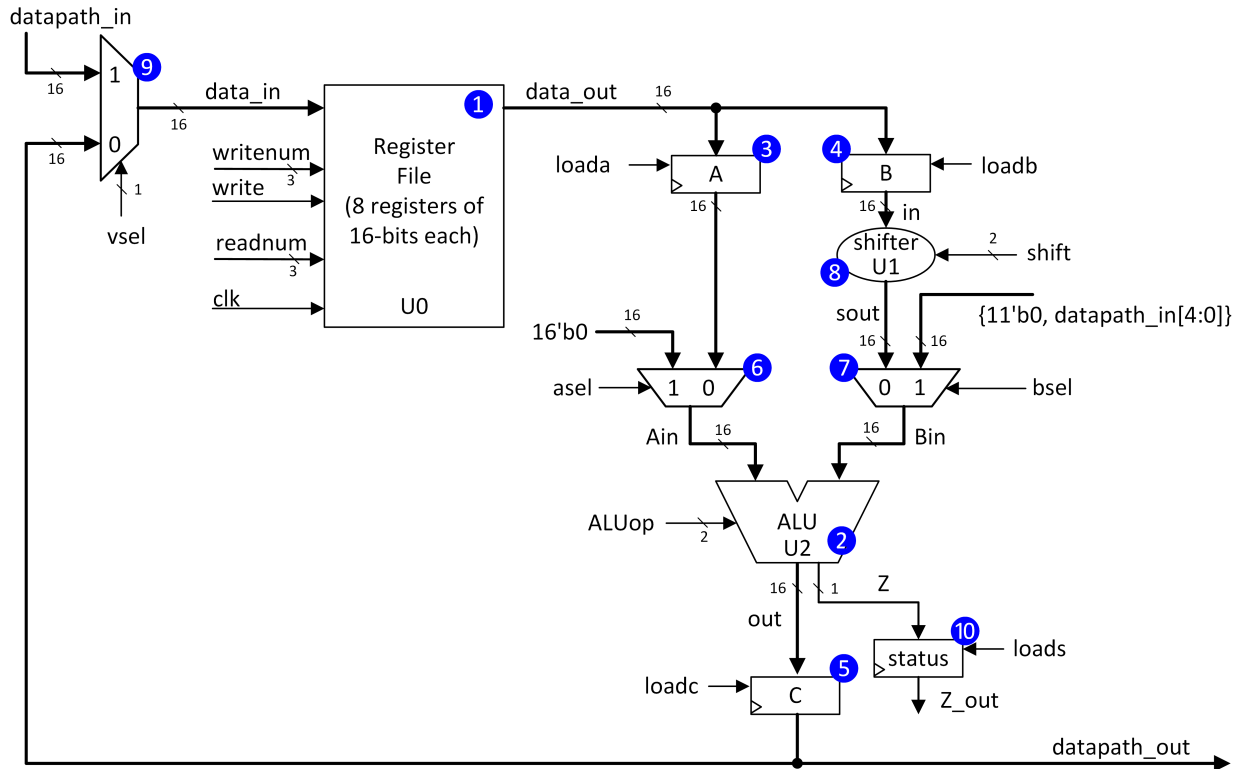


Figure 1: “Simple RISC Machine” Datapath




Figure 2: Register with Load Enable Circuit

value of out is updated to the value on in on the rising edge of the clock. In Figure 1 a 16-bit register with load enable is represented using the symbol shown in Figure 3.

Implement the register file out of eight instances of the 16-bit register with load enable circuit along with two binary to one-hot decoders, as shown in Figure 4. To be compatible with the autograder, your register file **must** contain signals named R0, R1, R2, R3, R4, R5, R6, R7 (Figure 4 omits R4, R5 and R6 for clarity).

Now, let’s consider how the value of variable “j” gets into and out of R3 inside the register file.

Suppose we want “j” to have the value 42 *before* we start executing our code. To put 42 in R3 you would place the 16-bit value 000000000101010 (binary for 42) on data_in, set the 3-bit input writenum to 011 (binary for 3), set write to 1 to indicate we wish to save, or *write*, the value 42 into location 3 in the register file, and input a rising edge on clk. This causes, the output of the upper 3:8 decoder to be driven to 00001000. Each output bit of the decoder is AND’ed with write. As write is 1, the load input to R3 is set to 1. On the rising edge of clk 000000000101010 will be copied to the 16-bit Q output of R3. At most one load enable input to R0 through R7 will be 1. If write is 0 all 8 load-enable signals are 0. Section 2 describes in more detail the “MOV” instruction that is used to write a constant into a register using the above steps.




Figure 3: Register with Load Enable Symbol (Note: this symbol is used in both Figure 1 and 4)

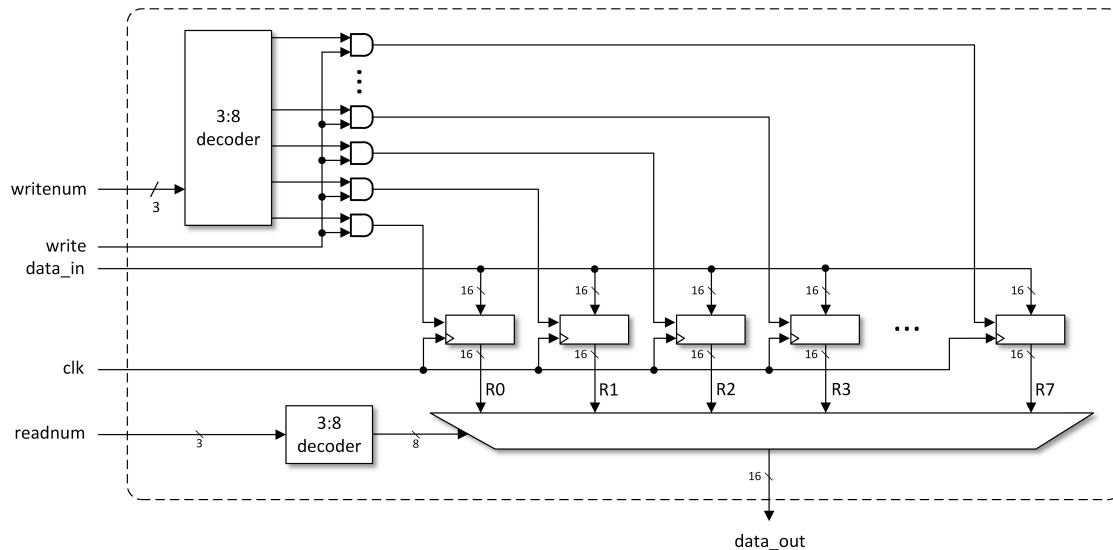



Figure 4: Register File Internal Structure (must contain signals R0, R1, ... R7 connected as shown)

To *read*, the value of “j”, which is in R3, we set the 3-bit bus *readnum* to 011. The 8-bit output of the lower 3:8 decoder will be 00001000 and this will cause the 8-input one-hot select mux to copy the value of R3 to *data_out*.

This register file is said to have two *ports*: one *write* port and one *read* port. Thus, up to one write and one read can be performed simultaneously. The read is combinational: whenever *readnum* changes, the value from the indicated register is driven out of the register file after some combinational delay. **THE REGISTER READS ARE NOT COORDINATED TO THE CLOCK!** The register write, however, is coordinated to the clock. At each rising clock edge, if *write* is 1, the value on the 16-bit register file input *data_in* is written into the register indicated by the value on *writenum*. This write only happens on the rising clock edge. If, at the clock edge, *write* is 0, no register is updated. Be sure you follow the Verilog style guidelines for your register file code.

1.1.2 Arithmetic Logic Unit

The  symbol, labeled 2, represents an *arithmetic logic unit* (ALU). The ALU can perform arithmetic or logical operations. This is the hardware that actually “computes” things inside a computer. The operation to be performed is indicated by the value of *ALUOp* as shown in Table 1.

Your ALU should be purely combinational (there is no clock input). Whenever one of the inputs or *ALUOp* lines change, the output changes appropriately (the Verilog “+” and “-” operations are combinational). If the main 16-bit result of the ALU (*out*) is zero the 1-bit Z output should be 1 and otherwise 0.

Value on ALUop input	Operation
00	Ain + Bin
01	Ain - Bin
10	Ain & Bin
11	~Bin

Table 1: ALU operations

shift	Operation
00	B
01	B shifted left 1-bit, least significant bit is zero
10	B shifted right 1-bit, most significant bit, MSB, is 0
11	B shifted right 1-bit, MSB is copy of B[15]

Table 2: Shift Operation Encoding

This output will be used in later labs for compare and branch instructions.

1.1.3 Pipeline Registers

The datapath contains three special 16-bit registers with load enable 345. These hold the datapath signals A, B, and C. We will use these registers while executing an individual instruction. We need at least one of the two registers A and B because the ALU is purely combinational and the register file can read out only one of R0 through R7 at a time. You may want to try to eliminate the other two registers in Lab 8 but for now you should keep them.

1.1.4 Source Operand Multiplexers

To enable more complex instructions besides addition and subtraction it is helpful if we can change the inputs to the ALU via the source operand multiplexers 67. For some instructions added in later labs these multiplexers are used to set the 16-bit Ain input to the ALU to zero. For other instructions we use them to input an “immediate operand”.

1.1.5 Shifter Unit

Inspired by the ARMv7 instruction set, the Simple RISC Machine allows one source operand of an instruction can be shifted before an operation for certain instructions. To enable this the shifter unit 8 is a purely combinational logic block that operates as follows. The shifter takes one 16-bit input from the Q output of register B 4 and outputs either the same value or the value shifted one bit to the left or right according to the value on “shift” as described in Table 2.

For example, if the input to the shifter was 1111000011001111, then the output of the shifter would be as shown in Table 3.

shift	Output of shifter
00	1111000011001111
01	1110000110011110
10	0111100001100111
11	1111100001100111

Table 3: Example shift operations starting with 1111000011001111

1.1.6 Writeback Multiplexer

Once the ALU has computed a value the main 16-bit result is captured in register C. If we want to use this value as the input to a subsequent instruction we need to write it into the register file. However, we will also want to input values into the register file from other sources. Thus, we add a writeback multiplexer 9.

1.1.7 Status Register

The status register 10 is used to remember if the output of a given ALU computation had a special outcome indicating which direction and if/else statement or loop should go. Specifically, the status register should record the Z output value from the ALU if loads is set to 1 and otherwise the status register should keep its current output the same.

2 Example Datapath Operation

Consider the addition of two registers, R2 and R3 with the result being stored in R5. The addition takes four clock cycles. During the first cycle, readnum is set to 2 to indicate we want to read the 16-bit contents of R2 from the register file. At the same time loada is set to 1 to indicate that register A should be updated on the next rising edge of the clock (note the little triangle in the bottom left of register A indicates a clock input). During the second cycle, we set loada back to 0, set readnum to 3 to indicate we now want to read the 16-bit contents of R3. At the same time loadb is set to 1. With these control input settings on the next rising edge of the clock the contents of R3 will be copied to register B. During the third cycle loadb is set back to 0, ALUop is set to “00” to indicate addition (see Table 1 above), asel is set to zero to ensure the value in register A appears at the Ain input to the ALU. Similarly, bsel is set to 0 to indicate the output of the shifter unit appears at the Bin input to the ALU. The shifter unit is combinational logic that takes the 16-bit contents of B as its input and outputs the value either unmodified, or shifted to the left or right by one bit position depending upon the control input “shift”. During this cycle the shift input is set to “00” to indicate the value in B should not be shifted. Also during this cycle, loadc is set to 1 to ensure the result of the addition is saved in register C on the next rising edge of the clock. We can optionally set loads to 1 if we want to record the “status” of the computation. In this lab the status will simply indicate if the 16-bit result of the ALU was zero. In later labs we will see how this status information can be used to help implement “if” statements and “for” loops in a language like C or Java. During the fourth cycle, loadc is set back to 0, vsel is set to 0, write is set to 1, and writenum is set to 5. Together these cause the value in register C to be fed back and written into register R5 within the register file.

Note that during the fourth cycle, the value fed back also appears on the output pins datapath_out. For this lab you can connect datapath_out to the 7-segment displays on the DE1-SoC using the logic provided in lab5_top.sv. This will be the primary way to tell if your datapath is working when it is on the DE1-SoC. However, this is not the fastest or easiest way to debug your circuit.

As alluded to earlier, the basic interface between hardware and software inside a computer is through “instructions”. Each instruction tells the computer how to move and operate upon some data. Consider an instruction of the form “MOV R2, #32”. This instruction would load the actual number 32 into register R2. This can be performed with the datapath as follows:

During the first cycle, assume the number 32 appears on the 8-bit signal datapath_in (in a later lab, we will consider more realistic memory read/write strategies). During this same cycle, vsel is set to 1, write is set to 1, and the number 2 (indicating register #2) is driven on writenum. Note that this instruction can be performed in only one cycle, unlike the ADD instruction, which takes 4 cycles. This will become important in the next assignment.

If you have trouble following the above discussion, please also review the slides “Lab 5 Introduction” on Piazza. Also, the following video shows your instructor demonstrating a working Lab 5 implementation using a DE1-SoC showing how executing instructions will look when you are done: https://cpen211.ece.ubc.ca/video/lab5_demo/lab5_demo.html.

3 Lab Procedure

You will get through the lab far more quickly if you break down the overall work into smaller parts and complete, compile and test (using a testbench) each one *before* moving on to the next. If you are working with a partner, remember you must use “pair programming”, which means you should not divide up the work but rather meet and work together on a single solution. You can work together remotely using MS Visual Code Live Share. If you are worried you will not have time to do the entire lab then see the marking guideline in Section 4 to see how you can earn part marks.

3.1 Recommended Development Sequence

1. Create a new ModelSim project called lab5.
2. Add a file `regfile.sv` and write synthesizable code for your register file in this file. Make sure your register file uses exactly the following declaration which is required by the autograder (note Verilog is case sensitive):

```
module regfile(data_in,writenum,write,readnum,clk,data_out);
    input [15:0] data_in;
    input [2:0] writenum, readnum;
    input write, clk;
    output [15:0] data_out;
    // fill out the rest
endmodule
```

3. Add a file `regfile_tb.sv` to your project for your register file testbench module and inside it include your testbench module `regfile_tb` making sure to use lower case for the module name as this is required by the autograder. Your `regfile_tb` **must** contain a signal `err` that is initialized to 0 and set to 1 if an error is found and stays at 1 thereafter. The script part of your `regfile_tb` testbench should finish all tests before time 500. Section 8 provides some tips on how to write good unit level testbenches and introduces some Verilog syntax for automatically checking if the output results are correct to enable what is known as regression testing.
4. Compile `regfile.sv` in Quartus to check for synthesis errors (inferred latches, etc.). Check the resulting hardware that Quartus generates using “Tools” > “Netlist Viewers” > “RTL Viewer” to verify the hardware looks as you expect (e.g., combinational logic or flip-flops). Then, test and debug your design by compiling and simulating `regfile_tb.sv` along with `regfile.sv` in ModelSim. Use RTL level simulation so that your testbench can check the values of R0, R1, etc.
5. After you have debugged the register file should you go through the above steps for the ALU. Use the file names `alu.sv` and `alu_tb.sv`, ensure your testbench module is named `ALU_tb` and ensure your ALU module has the following declaration required by the autograder (note Verilog is case sensitive):

```
module ALU(Ain,Bin,ALUop,out,Z);
    input [15:0] Ain, Bin;
    input [1:0] ALUop;
    output [15:0] out;
    output Z;
    // fill out the rest
endmodule
```

Your `ALU_tb` must contain a signal `err` that is initialized to 0 and set to 1 if an error is found and stays at 1 thereafter. The script part of your `ALU_tb` testbench should finish all tests before time 500.

6. Implement your shifter module in `shifter.sv` and the testbench for your shifter in `shifter_tb.sv`. Ensure your testbench module is named `shifter_tb` and ensure your shifter module has the following declaration required by the autograder (note Verilog is case sensitive):

```
module shifter(in,shift,sout);
    input [15:0] in;
    input [1:0] shift;
    output [15:0] sout;
    // fill out the rest
endmodule
```

Your `shifter_tb` must contain a signal `err` that is initialized to 0 and set to 1 if an error is found and stays at 1 thereafter. The script part of your `shifter_tb` testbench should finish all tests before time 500.

7. Now that all three main datapath modules are trusted to work, instantiate them in your datapath and add the remaining building blocks. Instantiate each of the three units (Register file ①, ALU ② and Shifter ⑧) inside `datapath.sv`. Note that the autograder will assume your register file has the instance label `REGFILE` (in all caps) and that the input and outputs are consistent with the way they are referenced in `lab5_top.sv` and `lab5_autograder_check.sv` provided in the starter repo. Note the autograder and `lab5_top.sv` require that `asel`, `bsel` and `vsel` are binary select inputs. Next, add in the remaining logic blocks ③④⑤⑥⑦⑨⑩ to your datapath module using synthesizable Verilog that conforms to the style guidelines. Use no fewer than one always block or assign statement per hardware block in Figure 1. Register A, B, and C will each require an instantiated flip-flop module and an assign statement for the enable input in order to conform to the style guidelines.
8. Write a top level testbench module for your datapath called `datapath_tb` in `datapath_tb.sv`. To be compatible with the autograder your `datapath_tb` module must instantiate `datapath` using named port association (not implicit/positional) and it must define an internal signal called `err` that is initialized to 0 and set to 1 if an error is found and stays 1 thereafter. At a minimum your `datapath_tb` should test at least the sequence shown below:

```
MOV R0, #7           ; this means, take the absolute number 7 and store it in R0
MOV R1, #2           ; this means, take the absolute number 2 and store it in R1
ADD R2, R1, R0, LSL#1 ; this means R2 = R1 + (R0 shifted left by 1) = 2+14=16
```

However, to get full marks on the autograder your `datapath_tb` must be capable of identifying bugs in several alternative `datapath` modules we have crafted that containing design errors, and these may require more extensive test cases than required to get your own datapath bug free.

9. Test the overall datapath in ModelSim. If you see any suspicious outputs you should follow the debugging procedure in (b) to find relevant internal signals to add to the waveform viewer and restart and rerun the simulation.
10. Only after your overall design is working in ModelSim should you compile your top level and attempt to download to your DE1-SoC. Use `lab5_top.sv` ONLY to help with this step. If you encounter bugs here try step (d). If you still are not sure what is going on and why the results differ from ModelSim then modify `lab5_top.sv` to connect internal signals within your datapath to the LEDs on your DE1-SoC to help you follow the debugging rule “Quit Thinking and Look”.

4 Marking Scheme

A reminder that *both* partners must be in attendance during the demo. Any partner who is absent will automatically receive a mark of zero even if they did their fair share of the work.

The file `lab5_autograder_check.sv` is provided in the starter repo to enable you to verify your modules are defined consistently with what the autograder expects. **WARNING:** *The file `lab5_autograder_check.sv` is NOT the autograder that will be used mark you. Passing the checks in this file does NOT (in any way) guarantee you will not lose marks when your code is run through the actual autograder. You are responsible for designing your own test benches to verify you match the specification given earlier.* The file `lab5_autograder_check.sv` contains three test benches that you should run: The module `lab5_check_1` checks that your register file, shifter and ALU can be checked by the autograder. It should print out:

```
# CHECK #1 DONE: Your register file, shifter and ALU appear compatible with the
# autograder. ** NOTE ** You must also manually verify you had no simulation
# warnings (above) and that you have no inferred latches (e.g., using Quartus).'
```

The module `lab5_check_2` runs your unit level test benches. It should print out:

```
# CHECK #2 DONE: Your unit level testbenches appear compatible with the autograder.
# ** NOTE ** You must manually verify you had no simulation warnings
# by looking above this line in the transcript window in ModelSim.
```

The module `lab5_check_3` runs your datapath testbench. It should print out:

```
# CHECK #3 DONE: Your datapath testbench appears compatible with the autograder.
# ** NOTE ** You must manually verify you had no simulation warnings
# by looking above this line in the transcript window in ModelSim.
```

The module `lab5_check_4` checks that your datapath can be checked by the autograder. It should print out:

```
# CHECK #4 DONE: Your datapath appears to be compatible with the autograder.
# ** NOTE ** You must manually verify you had no simulation warnings
# (above) and that you have no inferred latches (e.g., using Quartus).
```

As noted in the messages above, it is important to verify you did not have any simulation warnings after simulating each of these test benches.

IMPORTANT: Check your submission repo carefully as you will lose marks if your submission does not contain a Quartus Project File (.qpf) and the associated Quartus Settings File (.qsf) that indicates which Verilog files are part of your project. This .qsf file is created by Quartus when you create a project. It is typically named `<top_level_module_name>.qsf` and contains (among others) lines indicating which Verilog files are to be synthesized. If you open up this .qsf file you should see lines that look like the following. The key part is that these line contain “VERILOG_FILE”:

```
set_global_assignment -name VERILOG_FILE shifter.sv
set_global_assignment -name VERILOG_FILE regfile.sv
set_global_assignment -name VERILOG_FILE lab5_top.sv
set_global_assignment -name VERILOG_FILE datapath.sv
set_global_assignment -name VERILOG_FILE alu.sv
```

The autograder will use your .qsf file to determine which Verilog files should be synthesized together. To be sure, note the above .qsf file is **not** the file `DE1_SoC.qsf` we provide in the starter repo for importing DE1-SoC pin assignments. Also remember to include a Modelsim Project File (.mpf), which **MUST** be called “`lab5_top.mpf`” and must include both your synthesizable verilog and your unit level testbench files (if you forget this file or it does not include your testbenches, you will at almost certainly get zero marks for Part 2 and/or Part 3). Finally include your programming (.sof) file and any waveform (.do) files. **You will lose one mark for each one of these files that is missing.**

Your mark will be computed partly by running it through an autograder and partly assigned by your TA as follows.

4.1 Autograder Marking [5 marks]

The autograder will assign a mark out of 5 as follows:

1.5 Marks 0.5 marks for each of `regfile.sv`, `alu.sv` and `shifter.sv`. For the autograder to not get confused, there must only be one file in your submitted .zip file with each of these names. Make sure your register file, ALU and shifter are free of inferred latches. Your register file, ALU and shifter must be defined in modules named `regfile`, `ALU` and `shifter` following the module declarations described in Section 3.1. You will get 0.5 marks for each unit that synthesizes, is free of inferred latches and/or other combinational-loops and passes the autograders tests.

1.5 marks The autograder will check your unit-level test benches which must be in `regfile_tb.sv`, `alu_tb.sv` and `shifter_tb.sv` and have module names `regfile_tb`, `ALU_tb` and `shifter_tb`. Each of `regfile_tb`, `ALU_tb` and `shifter_tb` must contain a signal `err` that is initialized to 0 then set to 1 if an error is found and stays at 1 thereafter. In addition, since the autograder will use designs for the ALU, register file and shifter that are different from your own your test benches should only check output signals. For each unit level test bench the autograder will try various incorrect designs to see how many your testbench flags as incorrect by setting `err` to 1. To receive full marks your testbench should catch “most” of our broken designs. As above, for the autograder to not get confused, there must only be one file in your submitted .zip file with each of these names. To be compatible with the autograder your unit level test benches should pass `lab5_check_1` and in addition **NOT** make any assumptions about the signal names inside `regfile`, `ALU` and `shifter` (i.e., do not use hierarchical signal names to refer to them) or you will get zero for this part. The sole exception is that you may access internal signals `R0`, `R1`, ... `R7` inside `regfile` since these names are required for the autograder for `datapath` as indicated in `lab5_check_4`.

1 mark For your top level testbench `datapath_tb` in `datapath_tb.sv`. Similar to the unit level test benches we will try several broken `datapath` modules to see if you catch our errors. As above, for the autograder to not get confused, there must only be one file in your submitted .zip file with this name. Your `datapath_tb` must contain a signal `err` that is initialized to 0 then set to 1 if an error is found and stays at 1 thereafter. The autograder will try various incorrect `datapath` designs to see how many your testbench flags as incorrect by setting `err` to 1. To receive full marks your testbench should catch “most” of our broken designs. To be compatible with the autograder your `datapath_tb` must pass `lab5_check_3` and in addition **NOT** make any assumptions about the signal names inside `datapath` or any modules it instantiates (i.e., do not use hierarchical signal names to refer to them) or you will get zero for this part. The sole exception is that you may access internal signals `R0`, `R1`, ... `R7` inside `regfile` since these names are required for the autograder for `datapath` as indicated in `lab5_check_4`.

1 mark For your `datapath` module which must be defined inside of `datapath.sv`. To get this mark your “.zip” file must contain the Quartus “.qsf” file containing `VERILOG_FILE` for each file required to synthesize your code as the autograder will use this file to find any files in your “.zip” required to synthesize your `datapath` module. Your `datapath` must also use the input output port names and widths implied by `lab5_top.sv` and `lab5_autograder_check.sv`. You will get this mark provided your complete design synthesizes without errors, is completely free of inferred latches and/or other combinational-loops (including in your ALU, shifter and register file) and passes our test benches.

4.2 Teaching Assistant assigned [5 marks]

Partners may get a different mark based upon their ability to answer the TA’s questions. If it becomes apparently one partner did more than two thirds of the work the partner who did less will receive a mark of zero. You **must** include at least one line of comments per always block, assign statement or module instantiation and in test benches you must include one line of comments per test case saying what the test is for and what the expected outcome is.

1.5 Marks For explaining to the TA your `regfile.sv`, `alu.sv`, `shifter.sv` and `datapath.sv` code Up to 1 mark may be deducted for violations of the style guidelines and/or for lack of comments.

1.5 Marks For explaining your test strategy in your test benches and showing your simulation waveforms. You may lose up to 1 marks if your test cases are not commented (one line per test case saying what is being tested and the expected outcome).

2 Marks For demonstrating your datapath works on your DE1-SoC using a test case of your own devising. Your TA will need to be convinced your design really works to get full marks. The following video should give you an idea how this part might look: https://courses.ece.ubc.ca/cpen211/2016/lab5_demo/lab5_demo.html. To ensure all students can be marked within the duration of your lab session your TAs will give you no more than 5 minutes to complete this part and may give you zero if you cannot complete the demo in this time. Given we will be marking you on Zoom one partner should explain what each step is while the other performs that step using their DE1-SoC (decide before your demo who will explain and who will perform the steps on their DE1-SoC).

5 Lab Submission

If you are working with a partner, your submission **MUST** include a file called “CONTRIBUTIONS.txt” that describes each student’s contributions to each file that was added or modified. If either partner contributed less than one third to the solution (e.g., in lines of code), you must state this in your CONTRIBUTIONS file and verbally inform the TA. Your TA will deduct 3 marks if CONTRIBUTIONS is missing and may deduct up to this amount if it lacks in meaningful detail. Note that submitted files may be stored on servers outside of Canada. Thus, you may omit personal information (e.g., your name, SN) from your files and refer to “Partner 1” and “Partner 2” in CONTRIBUTIONS.

In addition, if you used any artificial intelligence tools, such as ChatGPT or github copilot you **MUST** document all prompts or other inputs in a file “AI.txt”. There should be enough detail in AI.txt to generate either the same (or very similar) code in case of any allegation that your (System)Verilog was shared with other groups. A reminder you are not permitted to share AI prompts (sharing AI prompts will be treated as misconduct).

Submitted files may be stored on computer servers outside of Canada. Thus, you may omit personal information.

Submit your code using github classroom using the invite link for your section of Lab 3 available at https://cpen211.ece.ubc.ca/cwl/github_info.php. After you have submitted your changes you should see the changes on github.com in the repo created using the invite link.

5.1 Submission checklist

Ensure your repo has the following files:

6 Lab Demonstration Procedure

Lookup your assigned TA, their Zoom meeting link, your marking time and your place in the marking order at https://cpen211.ece.ubc.ca/cwl/ta_lookup.php. As in Lab 3 and 4, your TA will have the files you submitted via “handin”. However, ensure you have your DE1-SoC with you.

7 Suggestions and Debugging Tips

See the debugging video <https://youtu.be/2c3CZouKJKs> for details of how to trace a bug to its source. Some other tips and tricks that you may find helpful:

1. Code in an iterative fashion, meaning write a portion of the code and test it and correct bugs then write more code.

2. Debugging. In the very likely case that a signal (wire or reg) appears wrong in the waveform viewer, first find the Verilog corresponding to the hardware block that “drives” that signal. If there is an obvious error in the code for that block that can explain the exact wrong result you are seeing, then try fixing it. If there is no obvious error then you should not change the Verilog for that block! If you do make a change, remember to recompile your Verilog, restart the simulation and rerun the simulation. If your change did not fix the specific bug you were trying to fix, then undo it! This is important! If you make changes to your code that do not fix the bug they tend to make it harder to find the bug you were original interested in because the bug tends to “move around”. Instead of “undo” you can also comment out the “fix” code you added so you can get it back quickly in case you do end up needing it. Now, if/when you run out of things that could be wrong with the block defining the signal that looks wrong, do your best to guess which inputs to that block could lead to this incorrect output. Then add all the input signals to the block to the waveform viewer and restart the simulation and rerun it. If one of those inputs seems wrong, repeat Step 5 starting with the block that drives that signal.
3. If your code appears to work in simulation but not on the DE1-SoC, first check for and fix any inferred latch warnings found during synthesis. If this does not fix the problem, create a top-level RTL-level testbench that models the actions you are attempting using the switches on the board with the code in `lab5_top_tb.sv` and simulate it in ModelSim while adding internal signals. If the design appears to work pre-synthesis, then you may need to do post-synthesis simulations and temporarily add additional I/Os to your modules to check their values after synthesis (remember to remove those I/Os before submission).
4. SystemVerilog assertions. If you save your testbench file with the extension `.sv` and set the file properties to SystemVerilog you can make your testbench “self checking” by using the SystemVerilog assert statement. If we expect “s” to be `3'b100` at some point during the test script then we could write:

```
assert (datapath_tb.DUT.MUX1.s == 3'b100) $display("PASS");
else $error("FAIL");
```

If you want simulation to stop on an error go to “Simulate > Runtime Options...” then select the “Message Severity” tab and change the setting for “Break Severity” to “Error”.

5. The following Verilog “force” and “release” syntax can be helpful for debugging after you put your datapath together if you later find a bug. In ModelSim from a test script and using the above external name syntax you can override the logic value generated by the circuit itself to “inject” your own test values using the Verilog keyword “force”. Continuing the example above, suppose “s” has the value “010” but you would like to see what the output “b” of instance “m” is if instead “s” was “100”. You could write the following line in your Verilog testbench script to find out:

```
force datapath_tb.DUT.MUX1.s = 3'b100;
```

Later in your test script you can go back to using the value generated by the circuit by using “release”:

```
release datapath_tb.DUT.MUX1.s;
```

6. Revision Control, Incremental Testing, Regression Tests.

To make the best use of git and github make a habit of making frequent “commits” after every major change or improvement to the code. Revision control (git) can go “hand in hand” with a practice known as “regression testing”. The idea of regression testing is that any time you change something you re-run all testbenches you have written so far to ensure no test case that *was* passing stopped working.

Regression testing works best if you make your testbenches “self checking” by printing out a “PASS/-FAIL” message at the end of the simulation. Every time you make a change to your design (e.g., the register file or ALU) or test bench, you rerun ALL of your tests and only commit changes if all existing tests “pass”. To speed this up, run ModelSim from the command line. For example, for your Lab 3 code could create a new file “regress.do” containing:

```
vlib work
vlog lab3_top.sv
vlog tb_lab3.sv
vsim -c work.tb_lab3
run -all
quit -f
```

Then, type “vsim -c -do regress.do” and hit enter. You should see outputs like you would in the transcript window in ModelSim.

8 The Simple RISC Machine Instruction Set Architecture

The information in Table 4 and 5 is only relevant to Lab 6 and 7. An assembler will be provided to you for Lab 7. Each row in these tables specifies a single instruction. The assembly syntax is in the leftmost column. Each instruction is encoded using 16-bits. The next 16 columns indicate the binary encoding for the instruction. The last column on the right summarizes the operation of the instruction. The most significant 3-bits of each instruction (bits 15 through 13) are the opcode which indicates which instruction or class of instruction is represented.

Terminology quick definitions. These will be explained in more detail in the Lab 6 to 8 handouts.

- R_n, R_d, R_m are 3-bit register number specifiers.
- $im8$ is an 8-bit immediate operand encoded as part of the instruction.
- $im5$ is a 5-bit immediate operand encoded as part of the instruction.
- sh is a 2-bit immediate encoded as part of the instruction used to control the shifter. Legal values for $\langle sh_op \rangle$ are “LSL#1”, “LSR#1”, or “ASR#1”.
- $SX(f)$ sign extends the immediate value f to 16-bits.
- sh_R_m is the value of R_m after passing through the shifter connected to the Bin input to the ALU.
- $Z, V,$ and N are the zero, overflow and negative flags of the status register (only Z is implemented in Lab 5).
- **status** refers to all three of Z, V and N .
- $R[x]$ refers to the 16-bit value stored in register x .
- $M[x]$ is the 16-bit value stored in main memory (added in Lab 6) at address x .
- **PC** refers to the program counter register (added in Lab 6).
- $\langle label \rangle$ refers to a textual marker in the assembly that indicates an instruction address

Assembly Syntax (see text)	“Simple RISC Machine” 16-bit encoding																Operation (see text)
	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
Move Instructions	<i>opcode</i>			<i>op</i>		<i>3b</i>			<i>8b</i>								
MOV Rn, #<im8>	1	1	0	1	0	Rn			im8						R[Rn] = sx(im8)		
MOV Rd, Rm{, <sh_op>}	1	1	0	0	0	0	0	0	Rd	sh	Rm			R[Rd] = sh_Rm			
ALU Instructions	<i>opcode</i>			<i>ALUop</i>		<i>3b</i>			<i>3b</i>	<i>2b</i>	<i>3b</i>						
ADD Rd, Rn, Rm{, <sh_op>}	1	0	1	0	0	Rn			Rd	sh	Rm			R[Rd]=R[Rn]+sh_Rm			
CMP Rn, Rm{, <sh_op>}	1	0	1	0	1	Rn			0	0	0	sh	Rm			status=f(R[Rn]-sh_Rm)	
AND Rd, Rn, Rm{, <sh_op>}	1	0	1	1	0	Rn			Rd	sh	Rm			R[Rd]=R[Rn]&sh_Rm			
MVN Rd, Rm{, <sh_op>}	1	0	1	1	1	0	0	0	Rd	sh	Rm			R[Rd]=~sh_Rm			
Memory Instructions	<i>opcode</i>			<i>ALUop</i>		<i>3b</i>			<i>3b</i>	<i>5b</i>							
LDR Rd, [Rn{, #<im5>}]	0	1	1	0	0	Rn			Rd	im5						R[Rd]=M[R[Rn]+sx(im5)]	
STR Rd, [Rn{, #<im5>}]	1	0	0	0	0	Rn			Rd	im5						M[R[Rn]+sx(im5)]=R[Rd]	

Table 4: Data Processing and Data Movement Instructions

Assembly Syntax (see text)	“Simple RISC Machine” 16-bit encoding																Operation (see text)
	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
Branch	<i>opcode</i>			<i>op</i>		<i>cond</i>			<i>8b</i>								
B <label>	0	0	1	0	0	0	0	0	im8								PC = PC+1+sx(im8)
BEQ <label>	0	0	1	0	0	0	0	1	im8								if Z = 1 then PC = PC+1+sx(im8) else PC = PC+1
BNE <label>	0	0	1	0	0	0	1	0	im8								if Z = 0 then PC = PC+1+sx(im8) else PC = PC+1
BLT <label>	0	0	1	0	0	0	1	1	im8								if N != V then PC = PC+1+sx(im8) else PC = PC+1
BLE <label>	0	0	1	0	0	1	0	0	im8								if N!=V or Z=1 then PC = PC+1+sx(im8) else PC = PC+1
Direct Call	<i>opcode</i>			<i>op</i>		<i>Rn</i>			<i>8b</i>								
BL <label>	0	1	0	1	1	1	1	1	im8								R7=PC+1; PC=PC+1+sx(im8)
Return	<i>opcode</i>			<i>op</i>		<i>unused</i>			<i>Rd</i>	<i>unused</i>							
BX Rd	0	1	0	0	0	0	0	0	Rd	0	0	0	0	0	0	0	PC=Rd
Indirect Call	<i>opcode</i>			<i>op</i>		<i>Rn</i>			<i>Rd</i>	<i>unused</i>							
BLX Rd	0	1	0	1	0	1	1	1	Rd	0	0	0	0	0	0	0	R7=PC+1; PC=Rd
Special Instructions	<i>opcode</i>			<i>not used</i>													
HALT	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	go to halt state

Table 5: Control Instructions